

Collapsing the dispatch ceiling: kernel fusion for browser-native WebGPU quantum-circuit simulation

Ahmet Barış Günaydın

Independent researcher · github.com/abgnydn/webgpu-q

Abstract

Sequential GPU workloads in the browser are throttled by a fixed per-dispatch cost: every WebGPU compute dispatch pays a driver/submission overhead that, for fine-grained work, dominates the actual arithmetic. We quantify this overhead for browser quantum-circuit simulation — a per-gate dispatch cost of $\alpha \approx 5\text{--}500\ \mu\text{s}$ on commodity hardware — and show that **kernel fusion**, executing a chain of k quantum gates inside a single WGSL `main()`, drives the effective per-gate cost toward $\alpha/k + C$. We implement fusion as hand-written dense-tile WGSL kernels (2-, 3-, and 4-qubit dense operators) and measure a tier ladder against the simulator’s own unfused multi-dispatch path: brick-wall two-qubit fusion reaches **3.04×**, three-qubit-tile (Tier-C) fusion **4.22×**, both at fidelity $F \geq 0.99999$ vs the unfused statevector. Four-qubit-tile (Tier-D) fusion is an **honest negative** — it tops out at **3.78×** because it crosses from the dispatch-bound into the compute-bound regime, where the per-block complex-multiply count grows faster than the memory traffic fusion saves. Fusion is correctness-preserving — across 360 (N, k, trial) configurations the worst fidelity loss against the unfused statevector is $1 - F = 2.4 \times 10^{-6}$ — and pushes the fused statevector update toward memory-bandwidth-bound throughput. All results run in a browser tab with no installation; every number is backed by a committed, environment-stamped JSON artifact. We do not benchmark against native GPU simulators (cuQuantum [1], Qiskit Aer [2]) — those require hardware unavailable to us — and position relative to them only through the bandwidth-fraction argument.

1 Introduction

A WebGPU compute dispatch is not free. Beyond the shader’s own runtime, each `dispatchWorkgroups` call pays a driver-side submission and scheduling cost. For coarse-grained work this is negligible; for *fine-grained, sequential* work — where the natural unit of computation is small and many units run in strict order — it becomes the dominant term. Quantum-circuit statevector simulation is a clean instance: each gate is a small $(2 \times 2, 4 \times 4)$ operator applied to a large amplitude array, and gates within a circuit must be applied in order. A one-gate-per-dispatch scheduler spends more time in the driver than in the ALUs.

We measure this overhead directly, establish the per-gate cost as a function of circuit width and fusion factor, and show that fusing gate chains into single dispatches collapses it. The contribution is not a new simulation algorithm — statevector simulation is textbook — but a quantification of the WebGPU dispatch ceiling for sequential workloads [3] and a fidelity-validated demonstration of how far kernel fusion moves it, including the regime where fusion stops paying off. Single-kernel fusion has previously removed this overhead for two other sequential WebGPU workloads — evolutionary fitness evaluation, at $159\text{--}720\times$ [4], and autoregressive transformer decoding, at $66\text{--}458\times$ [5] — both of which are deeply dispatch-bound. Quantum-circuit statevector simulation is a third domain, and, as we show, one that sits far closer to the bandwidth-bound end of that spectrum: each

gate already moves the entire 2^N -amplitude statevector, so the per-dispatch overhead is a much smaller fraction of the work than in those deeply dispatch-bound workloads, and dense-tile fusion additionally crosses into a compute-bound regime (§3.3). The same technique that delivers hundreds-fold speedups there therefore caps at a few-fold here. Locating where a workload sits on that dispatch-bound-to-bandwidth-bound spectrum is the point. The work runs entirely in a browser tab.

2 Methods

2.1 Simulator and fusion kernels

Amplitudes are stored as interleaved `vec2<f32>` (re, im) in a storage buffer of size 2^{N+3} bytes for N qubits. A single-qubit gate dispatches $N/2$ threads, each updating the amplitude pair that differs in the target bit; a controlled two-qubit gate dispatches $N/4$ threads. Fusion replaces a chain of such gates with one dense operator applied in a single dispatch: a brick-wall layer of adjacent two-qubit gates becomes one dense 4×4 dispatch (Tier-B); a three-qubit tile becomes a dense 8×8 dispatch (Tier-C); a four-qubit tile a dense 16×16 dispatch (Tier-D). The dense-tile operators are hand-written WGSL (`two-qubit-dense.wgsl`, `three-qubit-dense.wgsl`, `four-qubit-dense.wgsl`); the fused operator matrix is precomputed on the host and uploaded once per fused block. (This is fixed-tile fusion, not a general gate-chain JIT compiler.)

2.2 Measurement harness

Wall-clock is measured with `performance.now()` bracketed by a forced GPU sync (a mapped readback of a tiny buffer) before and after, because `queue.submit` is non-blocking and raw timing is otherwise fiction. Five warmup samples are discarded and twenty retained; we report the median and flag any configuration with `std/median > 0.1` as `noisy`. Correctness uses fidelity $F = |\langle \psi_{\text{ref}} | \psi_{\text{test}} \rangle|^2$ against the unfused statevector, with a pass bar of $F \geq 1 - 10^{-5}$ for f32 amplitude paths. Every run emits a JSON artifact recording git SHA, `adapter.info`, WebGPU limits, OS, seeds, warmup/trials, and the pass bar.

3 Results

All measurements on an M2 Pro under Chromium with WebGPU.

3.1 The dispatch ceiling and its collapse

The per-gate dispatch cost is $\alpha \approx 5\text{--}500\ \mu\text{s}$ on commodity WebGPU (level-1 dispatch-overhead experiment). Fusing k gates per dispatch drives the effective per-gate cost as $\alpha_{\text{eff}}(k) = \alpha/k + C$: measured single-qubit chains at $N = 8$ fall from $54.5\ \mu\text{s}/\text{gate}$ at $k = 1$ to $17.8\ \mu\text{s}$ at $k = 4$ and $\approx 15.8\ \mu\text{s}$ at $k = 8$, the $1/k$ signature of amortizing a fixed cost over k operations. (This experiment is flagged `noisy` — sub- $20\ \mu\text{s}$ timings sit near the harness’s sync-fence resolution — but the $1/k$ trend is unambiguous across widths.)

3.2 Correctness

Fusion is correctness-preserving. Across 360 (N, k , trial) cells the worst fidelity is $F = 0.99999762$ ($1 - F = 2.4 \times 10^{-6}$), inside the 10^{-5} bar; the worst state-norm deviation is 10^{-6} . Fusion changes *when* arithmetic happens, not *what* it computes.




Figure 1: The per-gate dispatch cost collapses as $1/k$. Measured effective per-gate cost for $N=8$ single-qubit chains (markers) against the $\alpha/k + C$ model (dashed). The cost falls toward a fixed asymptote $C \approx 14 \mu\text{s}/\text{gate}$ as the fusion factor k grows; the $k=1$ point sits above the asymptotic model because the first unfused dispatch carries additional fixed submission overhead.

3.3 The fusion tier ladder

tier	fusion	bound	best speedup	worst F	verdict
B (brick-wall)	2-qubit $\rightarrow$ dense 4×4	dispatch	3.04 $\times$	1.0000000	pass ($\geq 2\times$)
C (three-qubit tile)	5 ops $\rightarrow$ dense 8×8	dispatch	4.22 $\times$	0.9999988	pass ($\geq 3\times$)
D (four-qubit tile)	7 ops $\rightarrow$ dense 16×16	dispatch $\rightarrow$ compute	3.78 $\times$	0.9999995	fail ($\geq 5\times$)

Best speedups measured at: B — $N=20$, depth 80; C — $N=15$, depth 80; D — $N=20$, depth 80.

Tier-D is the informative result. Its correctness is fine (worst $F = 0.9999995$), but it tops out at $3.78\times$ against a $5\times$ target because the four-qubit dense block performs 256 complex multiplies per 16-amplitude tile (16 per amplitude) — a $4\times$ growth in per-tile arithmetic over Tier-C’s 64, against only a $2\times$ growth in the memory traffic that fusion eliminates. Tier-D crosses out of the dispatch-bound regime into the compute-bound one, where collapsing dispatches no longer helps. This locates the boundary of the technique rather than hiding it.

3.4 Toward bandwidth-bound throughput

In the dispatch-bound regime (small N , long circuits) fused same-qubit chains at $k = 32$ reach $2.62\times$ ($N = 14$, depth 160). It is tempting to read the fused path’s *logical* effective throughput — the statevector traffic of all D gates, $16 D 2^N/t$, which rises to 866 GB/s at $N = 20$, depth 160 — as memory-bandwidth saturation. It is not: that figure exceeds this device’s ≈ 200 GB/s DRAM bandwidth by $4.3\times$ precisely *because* fusion is working — a $k = 32$ chain reads each amplitude from DRAM once, applies all 32 gates in registers, and writes once, so it does the work of 866 GB/s of naive per-gate traffic while moving only one read+write pass per fused dispatch. The *physical* DRAM bandwidth, $16 \lceil D/k \rceil 2^N/t$, peaks at just ~ 27 GB/s — **13.5% of the device bound** — so the fused kernel is *not* bandwidth-saturated in this regime. The tell is in the artifact: the




Figure 2: Fusion tier ladder. Brick-wall (B) and three-qubit-tile (C) fusion clear their speedup targets; four-qubit-tile (D) fusion plateaus at $3.78\times$ against a $5\times$ target because the dense 16×16 block is compute-bound — 256 complex multiplies per 16-amplitude tile (16 per amplitude), a $4\times$ growth in per-tile arithmetic over Tier-C’s 8×8 block (64 per tile), while the memory traffic fusion eliminates grows only $2\times$. The plateau is committed as a `status:"fail"` artifact, not hidden.

fused wall-time is pinned at $\sim 3\text{ms}$ across $N = 10$ to $N = 20$ (a $1000\times$ range of state sizes), i.e. floor-limited by the harness sync fence rather than by memory traffic. The fused-vs-unfused speedup compresses toward $1.25\times$ at low depth ($N = 20$, depth 40) because few dispatches remain to remove, not because bandwidth is saturated; at longer depth the dispatch term still dominates and the gain persists ($N = 20$, depth 160 stays $\sim 2.6\times$). True DRAM saturation is the asymptotic ceiling fusion drives *toward* as N grows past the floor regime; demonstrating it empirically — at statevector sizes where the per-pass compute time dwarfs the $\sim 3\text{ms}$ fence — is future work (§4). All figures here are derived from the committed `E10-throughput` artifact’s (N, D, k, t) fields.

4 Honest limitations

- **Tier-D plateau (§3.3)** — fusion past three-qubit tiles is compute-bound and does not reach the $5\times$ target. Committed as a `status:"fail"` artifact with the compute-vs-memory diagnosis.
- **Bandwidth saturation is asserted asymptotically, not measured.** At the widths we test ($N \leq 20$) the fused wall-time is floor-limited by the harness sync fence ($\sim 3\text{ms}$, constant across a $1000\times$ range of state sizes), so physical DRAM utilization peaks at only $\sim 13.5\%$ of the device bound. We therefore do *not* claim the fused kernel saturates memory bandwidth in the measured regime; establishing the true asymptotic ceiling needs statevector sizes large enough that per-pass compute dwarfs the fence ($N \gtrsim 24$), a measurement we have not yet run.
- **No native-simulator head-to-head.** We measure fusion against this simulator’s own unfused dispatch path. We did *not* benchmark cuQuantum `custatevec` [1] or Qiskit Aer GPU [2] — both require NVIDIA hardware we do not have — and we make no head-to-head claim against them. Positioning is via the bandwidth-fraction argument only.

- **Fixed-tile fusion, not a JIT chain compiler.** The dense operators are hand-written for 2/3/4-qubit tiles; a general gate-chain fuser that emits WGS� per circuit is future work.
- **f32 amplitudes** — the 10^{-5} fidelity bar reflects single-precision GPU arithmetic; f64 is not available in WebGPU 1.0 [6].
- **No subgroup operations** — WebGPU subgroups (shuffle/reduce) are out of the 1.0 spec [6]; they would unlock a further $\sim 2\times$ on the reduction-heavy paths when available.

5 Related work

The per-dispatch overhead of WebGPU is an actively characterized quantity [3]. Native GPU statevector simulators — NVIDIA cuQuantum custatevec [1] and Qiskit Aer GPU [2] — are the performance references for the algorithm, and both are known to approach memory-bandwidth-bound throughput for long contiguous gate chains; we target the same bandwidth-fraction regime but in a browser, and do not benchmark against them directly (§4). Kernel/operator fusion to amortize launch overhead is standard in GPU machine-learning runtimes; the contribution here is its quantification and fidelity-validated application to ordered quantum-gate chains under the specific constraints of WebGPU [6] in a browser tab. This is the third domain in which we characterize single-kernel WebGPU fusion: prior work established the framework for evolutionary fitness evaluation [4] and autoregressive transformer decoding [5], both deeply dispatch-bound and so reaching 10^2 – $10^3\times$ speedups. The present domain is the instructive opposite end of the spectrum — each statevector gate already moves the full amplitude array, so the workload is far less dispatch-bound to begin with, and dense-tile fusion crosses into compute-bound at Tier-D (§3.3) — so fusion’s payoff is bounded at a few-fold, and that plateau marks the crossover explicitly.

6 Software availability and reproducibility

Source: github.com/abgnydn/webgpu-q (MIT). Fusion kernels in `src/shaders/` (`*-dense.wgsl`); experiments in `experiments/level-3-fusion/` (E8 correctness, E9 dispatch-collapse, E10 throughput, E11 Tier-B, E12 Tier-C, E13 Tier-D); the committed `experiments/results/<date>/level-3/` JSON artifacts are the source of every number here. No installation: the simulator and its benchmarks run in any WebGPU-capable Chromium. Each artifact records git SHA, `adapter.info`, WebGPU device limits, OS, seeds, and the pass bar; timing is sync-fenced; warmup samples are discarded; failing configurations are committed with a diagnosis rather than rerun.

Generative-AI disclosure. Portions of the software, documentation, and this manuscript were drafted with a large language model used as a coding and writing aid; all output was author-reviewed, correctness was enforced by the fidelity tests and committed artifacts, and every quantitative claim is traceable to a recorded experiment.

Statements. Sole author; no competing interests; no external funding. Archived at Zenodo, concept DOI 10.5281/zenodo.20494382 (resolves to the latest version).

7 Conclusion

WebGPU’s per-dispatch overhead caps fine-grained sequential GPU work in the browser. For quantum-circuit simulation it is $\alpha \approx 5$ – $500\ \mu\text{s}/\text{gate}$, and kernel fusion drives the effective cost as

$\alpha/k + C$: a tier ladder of dense-tile WGS� kernels delivers up to $4.22\times$ over the unfused path at fidelity ≥ 0.99999 , with a clean, documented plateau at four-qubit tiles where the workload turns compute-bound. The result is a fidelity-validated map of how far fusion moves the dispatch ceiling — and exactly where it stops — entirely inside a browser tab, with every number reproducible from a URL.

References

- [1] Harun Bayraktar, Ali Charara, David Clark, Saul Cohen, Timothy B. Costa, et al. cuQuantum SDK: A high-performance library for accelerating quantum science. In *2023 IEEE International Conference on Quantum Computing and Engineering (QCE)*, pages 1050–1061, 2023. DOI: 10.1109/QCE57702.2023.00119. Preprint arXiv:2308.01999.
- [2] Ali Javadi-Abhari, Matthew Treinish, Kevin Krsulich, Christopher J. Wood, Jake Lishman, Julien Gacon, Simon Martiel, Paul D. Nation, Lev S. Bishop, Andrew W. Cross, Blake R. Johnson, and Jay M. Gambetta. Quantum computing with Qiskit, 2024. DOI: 10.48550/arXiv.2405.08810. arXiv:2405.08810. Aer GPU statevector backend.
- [3] Anonymous. Characterizing WebGPU dispatch overhead for LLM inference across four GPU vendors, three backends, and three browsers. *arXiv preprint*, arXiv:2604.02344, 2026.
- [4] Ahmet Barış Günaydın. Single-kernel fusion for sequential fitness evaluation via WebGPU compute shaders, 2026. Concept DOI: 10.5281/zenodo.19331833.
- [5] Ahmet Barış Günaydın. Single-kernel fusion for autoregressive transformer decoding via WebGPU compute shaders, 2026. Concept DOI: 10.5281/zenodo.19344276.
- [6] W3C GPU for the Web Working Group. WebGPU. W3C Working Draft, <https://www.w3.org/TR/webgpu/>, 2025.